

EECS E6893 Big Data Analytics - Homework Assignment 1

Name: Qi Wang

UNI: qw2261

```
In [1]: import operator
import sys
from pyspark import SparkConf, SparkContext
import numpy as np
import matplotlib.pyplot as plt
from scipy import linalg
```

```
In [2]: # Macros.
MAX_ITER = 20
DATA_PATH = "gs://homework0_qi/HW1/data.txt"
C1_PATH = "gs://homework0_qi/HW1/c1.txt"
C2_PATH = "gs://homework0_qi/HW1/c2.txt"
NORM = 2
```

```
In [3]: # Spark settings
# conf = SparkConf()
sc = SparkContext.getOrCreate()
```

```
In [4]: # Load the data, cache this since we're accessing this each iteration
data = sc.textFile(DATA_PATH).map(
    lambda line: np.array([float(x) for x in line.split(' ')]))
    ).cache()
# Load the initial centroids c1, split into a list of np arrays
centroids1 = sc.textFile(C1_PATH).map(
    lambda line: np.array([float(x) for x in line.split(' ')]))
    ).collect()
# Load the initial centroids c2, split into a list of np arrays
centroids2 = sc.textFile(C2_PATH).map(
    lambda line: np.array([float(x) for x in line.split(' ')]))
    ).collect()
```

```
In [5]: # Helper functions.
def closest(p, centroids, norm):
    """
    Compute closest centroid for a given point.
    Args:
        p (numpy.ndarray): input point
        centroids (list): A list of centroids points
        norm (int): 1 or 2
    Returns:
        int: The index of closest centroid.
    """
    closest_c = min([(i, linalg.norm(p - c, norm))
                     for i, c in enumerate(centroids)],
                     key=operator.itemgetter(1))[0]
    return closest_c
```

```

In [6]: # K-means clustering
def kmeans(data, centroids, norm=2):
    """
    Conduct k-means clustering given data and centroid.
    This is the basic version of k-means, you might need more
    code to record cluster assignment to plot TSNE, and more
    data structure to record cost.
    Args:
        data (RDD): RDD of points
        centroids (list): A list of centroids points
        norm (int): 1 or 2
    Returns:
        RDD: assignment information of points, a RDD of (centroid, (point, 1))
        list: a list of centroids
        and define yourself...
    """
    # iterative k-means
    within_cluster_cost = []
    training_data = data.collect()

    for _ in range(MAX_ITER):
        iter_cost = 0

        combo_points = []
        # Transform each point to a combo of point, closest centroid, count=1
        # point -> (closest_centroid, (point, 1))
        for point in training_data:
            closest_centroid = closest(point, centroids, norm)
            combo_points.append((closest_centroid, (point, 1)))

        for each_combo in combo_points:
            single_part = (linalg.norm(each_combo[1][0] - centroids[each_combo[0]], norm))
            if norm == 2:
                single_part = single_part ** 2
            iter_cost += single_part

        print(iter_cost)
        within_cluster_cost.append(iter_cost)

    # Re-compute cluster center
    # For each cluster center (key), aggregate its values

```

```

# by summing up points and count
combo_points_rdd = sc.parallelize(combo_points)
combo_points_rdd = combo_points_rdd.reduceByKey(lambda v1, v2: (v1[0] + v2[0], v1[1] + v2[1]))

# Average the points for each centroid: divide sum of points by count
combo_points_rdd = combo_points_rdd.map(lambda v: (v[0], (v[1][0] / v[1][1], v[1][1])))

# Use collect() to turn RDD into list
updated_centroids = []
for each_centroid in combo_points_rdd.sortByKey().collect():
    updated_centroids.append(each_centroid[1][0])

centroids = updated_centroids

return_combo = sc.parallelize(combo_points)
return return_combo, centroids, within_cluster_cost

```

Question one

(1). Using L1 distance as similarity measurement

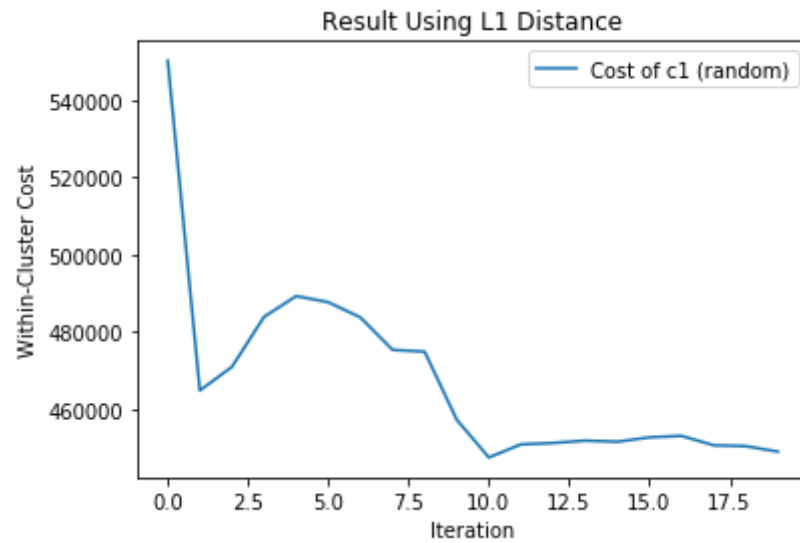
```
In [7]: L1_c1_combo, L1_c1_centroid, L1_c1_wc_cost = kmeans(data, centroids1, norm=1)
```

```
550117.1420000045  
464829.2684039464  
470934.15384668263  
483874.81628509297  
489234.2347883483  
487664.6926267901  
483718.66592851654  
475337.94763305597  
474871.9665496577  
457244.78974174923  
447493.195604051  
450891.8358047716  
451232.5774756949  
451860.12588546367  
451567.2235891512  
452710.0520999444  
453078.22696184984  
450646.13556209765  
450419.97011343326  
449009.59037188475
```


```
In [8]: L1_c2_combo, L1_c2_centroid, L1_c2_wc_cost = kmeans(data, centroids2, norm=1)
```

```
1433739.30999999938
1084488.7769648738
973431.7146620394
895934.5925630673
865128.3352940796
845846.6470313473
827219.5827561237
803590.3456011107
756039.5172761244
717332.9025432297
694587.9252526845
684444.5019967925
674574.7475478566
667409.469916026
663556.6278214998
660162.777228758
656041.3222947085
653036.7540731638
651112.4262522653
649689.0131843556
```


```
In [9]: x = range(20)
plt.plot(x, L1_c1_wc_cost)
plt.legend(['Cost of c1 (random)'])
plt.title("Result Using L1 Distance")
plt.xlabel('Iteration')
plt.ylabel('Within-Cluster Cost')
plt.show()
```



```
In [10]: x = range(20)
plt.plot(x, L1_c2_wc_cost)
plt.legend(['Cost of c2 (farthest initialization)'])
plt.title("Result Using L1 Distance")
plt.xlabel('Iteration')
plt.ylabel('Within-Cluster Cost')
plt.show()
```




```
In [11]: x = range(20)
plt.plot(x, L1_c1_wc_cost)
plt.plot(x, L1_c2_wc_cost)
plt.legend(['Cost of c1 (random)', 'Cost of c2 (farthest initialization)'])
plt.title("Result Using L1 Distance")
plt.xlabel('Iteration')
plt.ylabel('Within-Cluster Cost')
plt.show()
```



(2). Using L2 distance as similarity measurement


```
In [12]: L2_c1_combo, L2_c1_centroid, L2_c1_wc_cost = kmeans(data, centroids1, norm=2)
```

```
623660345.306  
509862908.298  
485480681.872  
463997011.685  
460969266.573  
460537847.983  
460313099.654  
460003523.889  
459570539.318  
459021103.342  
458490656.192  
457944232.588  
457558005.199  
457290136.352  
457050555.06  
456892235.615  
456703630.737  
456404203.019  
456177800.542  
455986871.027
```

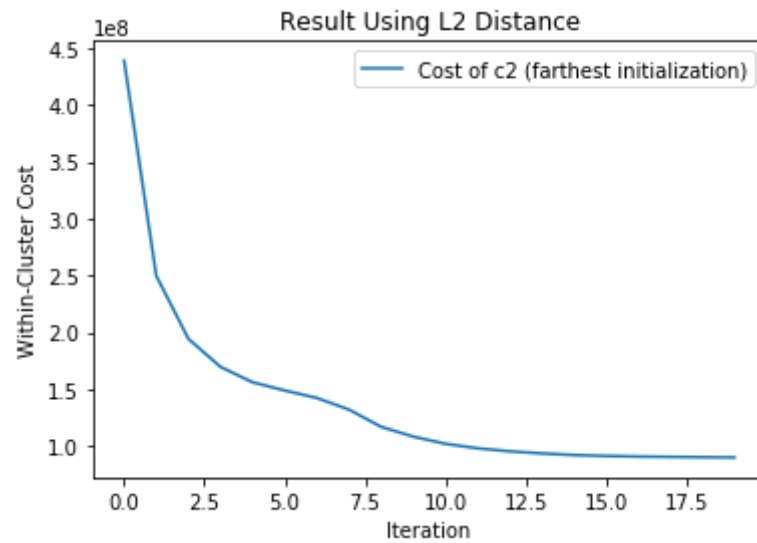
```
In [13]: L2_c2_combo, L2_c2_centroid, L2_c2_wc_cost = kmeans(data, centroids2, norm=2)
```

```
438747790.028  
249803933.626  
194494814.406  
169804841.452  
156295748.806  
149094208.109  
142508531.62  
132303869.407  
117170969.837  
108547377.179  
102237203.318  
98278015.7498  
95630226.1218  
93793314.0512  
92377131.9682  
91541606.2542  
91045573.8304  
90752240.1014  
90470170.1812  
90216416.1756
```


```
In [14]: x = range(20)
plt.plot(x, L2_c1_wc_cost)
plt.legend(['Cost of c1 (random)'])
plt.title("Result Using L2 Distance")
plt.xlabel('Iteration')
plt.ylabel('Within-Cluster Cost')
plt.show()
```



```
In [15]: x = range(20)
plt.plot(x, L2_c2_wc_cost)
plt.legend(['Cost of c2 (farthest initialization)'])
plt.title("Result Using L2 Distance")
plt.xlabel('Iteration')
plt.ylabel('Within-Cluster Cost')
plt.show()
```



```
In [16]: x = range(20)
plt.plot(x, L2_c1_wc_cost)
plt.plot(x, L2_c2_wc_cost)
plt.legend(['Cost of c1 (random)', 'Cost of c2 (farthest initialization)'])
plt.title("Result Using L2 Distance")
plt.xlabel('Iteration')
plt.ylabel('Within-Cluster Cost')
plt.show()
```



(3). T-SNE reduction and visualizing the clustering of result (2) in 2D space

```
In [17]: from sklearn.manifold import TSNE
```

```

In [18]: ## Help function to visualize the clustering situation
def T_SNE_graph(combo_points, random_seed, init_cluster):
    """
    Conduct T-SNE reduction for points
    to 2D space and visualize the clustering
    result in the graph.
    Args:
        combo_points (RDD): a RDD of (centroid, (point, 1))
    Returns:
        T_SNE_plot: a plot of clustering in 2D space
    """
    labels = np.array(combo_points.map(lambda v: v[0]).collect())
    points = np.array(combo_points.map(lambda v: v[1][0]).collect())
    points_embbed = TSNE(n_components = 2, random_state = random_seed).fit_transform(points)
    print('Shape: ' + str(points.shape) + ' -> ' + str(points_embbed.shape))

    plt.figure()
    plt.scatter(points_embbed[:, 0], points_embbed[:, 1])
    plt.title('%s Before Clustering' % init_cluster)

    T_SNE_plot = plt.figure()
    plt.scatter(points_embbed[:, 0], points_embbed[:, 1], c = labels)
    plt.title('%s After Clustering' % init_cluster)
    plt.show()
    return T_SNE_plot

```

```

In [19]: random_seed_given = np.random.randint(10)
         random_seed_given

```


```


Out[19]: 2

```

```
In [20]: t_sne_c1 = T_SNE_graph(L2_c1_combo, random_seed_given, 'Initial_Cluster: c1 (random cluster)')
```



Shape: (4601, 58) -> (4601, 2)





```
In [21]: t_sne_c2 = T_SNE_graph(L2_c2_combo, random_seed_given, 'Initial_Cluster: c2 (farthest initializationr)')
```

Shape: (4601, 58) -> (4601, 2)



(4). Determine performance of K-means initialization using c1.txt and c2.txt for L1 and L2

```
In [22]: (L1_c1_wc_cost[0] - L1_c1_wc_cost[-1]) / L1_c1_wc_cost[0]
```

```
Out[22]: 0.1837927668651324
```

```
In [23]: (L1_c2_wc_cost[0] - L1_c2_wc_cost[-1]) / L1_c2_wc_cost[0]
```

```
Out[23]: 0.5468569434813374
```

In terms of L1, c1 improves by 18.38% and c2 improves by 54.69%. However, c1 reaches the smaller cost than c2.

Based on my analysis, the metric of selecting centroids that are farthest is the better method, which means that c2 should be better than c1.

In this situation, the result displays that c1 is better than c2, which might result from the case that the "farthest" is based on Euclidean metrics rather than Manhattan metric and the "farthest" in Euclidean might not be real farthest in Manhattan.

If we need the real farthest, we should select farthest centroids based on Manhattan metric and do K-means again for comparison.

```
In [24]: (L2_c1_wc_cost[0] - L2_c1_wc_cost[-1]) / L2_c1_wc_cost[0]
```

```
Out[24]: 0.26885383292518295
```

```
In [25]: (L2_c2_wc_cost[0] - L2_c2_wc_cost[-1]) / L2_c2_wc_cost[0]
```

```
Out[25]: 0.7943775029159899
```

In terms of L2, c1 improves by 26.9% after 20 iterations and c2 improves by 79.4% after 20 iterations. Also, c2 reaches the smaller cost than c1.

Therefore, c2 is better than c1 in terms of L2 because it distributes the initial clusters far apart. There will be less overlap between points and the clusters are more obviously split, reaching a better result instance.

(5). The time complexity of iterative K-means

Assume a fixed number t of iterations, for n (d -dimensional) points, where k is the number of centroids (or clusters),
 $Time\ Complexity = O(t \times k \times n \times d) = O(tknd)$

Question two: Binary classification with Spark MLlib

(1). Read the csv and rename the columns

```
In [26]: import operator
         from pyspark.sql import SparkSession
```

```
In [27]: ## Spark settings
         spark = SparkSession.builder.appName('adult').getOrCreate()
```

```
In [28]: adult_PATH = 'gs://homework0_qi/HW1/adult.csv'
```

```
In [29]: ## Read the csv file
         df = spark.read.format("csv").options(header='false', inferschema='true').load(adult_PATH)
         type(df)
```

```
Out[29]: pyspark.sql.dataframe.DataFrame
```

```
In [30]: ## Rename the columns
original_columns = df.columns
updated_columns = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_status',
                   'occupation', 'relationship', 'race', 'sex', 'capital_gain', 'capital_loss',
                   'hours_per_week', 'native_country', 'income']

print('Original DataFrame is ', df)
print('')

for i in range(len(original_columns)):
    df = df.withColumnRenamed(original_columns[i], updated_columns[i])

print('Updated DataFrame is ', df)
print('')

('Original DataFrame is ', DataFrame[_c0: int, _c1: string, _c2: double, _c3: string, _c4: double, _c
5: string, _c6: string, _c7: string, _c8: string, _c9: string, _c10: double, _c11: double, _c12: doubl
e, _c13: string, _c14: string])

('Updated DataFrame is ', DataFrame[age: int, workclass: string, fnlwgt: double, education: string, ed
ucation_num: double, marital_status: string, occupation: string, relationship: string, race: string, s
ex: string, capital_gain: double, capital_loss: double, hours_per_week: double, native_country: strin
g, income: string])
```

(2). Data preprocessing

```
In [31]: from pyspark.ml.feature import OneHotEncoderEstimator
from pyspark.ml.feature import StringIndexer
from pyspark.ml.feature import VectorAssembler
```

```
In [32]: df.show(1)
```

```
+---+-----+-----+-----+-----+-----+-----+-----+-----+---+
+---+-----+-----+-----+-----+-----+-----+-----+-----+---+
|age| workclass| fnlwgt| education|education_num|marital_status| occupation| relationship| race|
sex|capital_gain|capital_loss|hours_per_week|native_country|income|
+---+-----+-----+-----+-----+-----+-----+-----+-----+---+
+---+-----+-----+-----+-----+-----+-----+-----+-----+---+
| 39| State-gov|77516.0| Bachelors|          13.0| Never-married| Adm-clerical| Not-in-family| White| M
ale|          2174.0|          0.0|          40.0| United-States| <=50K|
+---+-----+-----+-----+-----+-----+-----+-----+-----+---+
+---+-----+-----+-----+-----+-----+-----+-----+-----+---+
only showing top 1 row
```

```
In [33]: df.count()
```

```
Out[33]: 32561
```

```
In [34]: df.printSchema()
```

```
root
|-- age: integer (nullable = true)
|-- workclass: string (nullable = true)
|-- fnlwgt: double (nullable = true)
|-- education: string (nullable = true)
|-- education_num: double (nullable = true)
|-- marital_status: string (nullable = true)
|-- occupation: string (nullable = true)
|-- relationship: string (nullable = true)
|-- race: string (nullable = true)
|-- sex: string (nullable = true)
|-- capital_gain: double (nullable = true)
|-- capital_loss: double (nullable = true)
|-- hours_per_week: double (nullable = true)
|-- native_country: string (nullable = true)
|-- income: string (nullable = true)
```

```
In [35]: all_string_columns = ['workclass', 'education', 'marital_status', 'occupation',  
                             'relationship', 'race', 'sex', 'native_country', 'income']
```

```
In [36]: def dataPreprocessing(df, string_column):  
    index_output = string_column + 'Index'  
    vec_output = string_column + 'Vec'  
  
    indexer = StringIndexer(inputCol = string_column, outputCol = index_output)  
    df = indexer.fit(df).transform(df)  
  
    encoder = OneHotEncoderEstimator(inputCols = [index_output], outputCols = [vec_output])  
    df = encoder.fit(df).transform(df)  
  
    return df
```

```
In [37]: df_clean = df  
    for each_string_column in all_string_columns:  
        df_clean = dataPreprocessing(df_clean, each_string_column)
```

```
In [38]: df_clean.printSchema()
```

```
root
|-- age: integer (nullable = true)
|-- workclass: string (nullable = true)
|-- fnlwgt: double (nullable = true)
|-- education: string (nullable = true)
|-- education_num: double (nullable = true)
|-- marital_status: string (nullable = true)
|-- occupation: string (nullable = true)
|-- relationship: string (nullable = true)
|-- race: string (nullable = true)
|-- sex: string (nullable = true)
|-- capital_gain: double (nullable = true)
|-- capital_loss: double (nullable = true)
|-- hours_per_week: double (nullable = true)
|-- native_country: string (nullable = true)
|-- income: string (nullable = true)
|-- workclassIndex: double (nullable = false)
|-- workclassVec: vector (nullable = true)
|-- educationIndex: double (nullable = false)
|-- educationVec: vector (nullable = true)
|-- marital_statusIndex: double (nullable = false)
|-- marital_statusVec: vector (nullable = true)
|-- occupationIndex: double (nullable = false)
|-- occupationVec: vector (nullable = true)
|-- relationshipIndex: double (nullable = false)
|-- relationshipVec: vector (nullable = true)
|-- raceIndex: double (nullable = false)
|-- raceVec: vector (nullable = true)
|-- sexIndex: double (nullable = false)
|-- sexVec: vector (nullable = true)
|-- native_countryIndex: double (nullable = false)
|-- native_countryVec: vector (nullable = true)
|-- incomeIndex: double (nullable = false)
|-- incomeVec: vector (nullable = true)
```



```
In [39]: ## Assemble the vectors
input_col = ['age', 'workclassVec', 'fnlwgt', 'education_num', 'marital_statusVec',
             'occupationVec', 'relationshipVec', 'raceVec', 'sexVec', 'capital_gain',
             'capital_loss', 'hours_per_week', 'native_countryVec']
assembler = VectorAssembler(inputCols = input_col, outputCol = 'feature')

df_clean = assembler.transform(df_clean)
```

```
In [40]: df_clean.select(['feature', 'incomeIndex']).show()
```

feature	incomeIndex
(85,[0,5,9,10,12,...]	0.0
(85,[0,2,9,10,11,...]	0.0
(85,[0,1,9,10,13,...]	0.0
(85,[0,1,9,10,11,...]	0.0
(85,[0,1,9,10,11,...]	0.0
(85,[0,1,9,10,11,...]	0.0
(85,[0,1,9,10,16,...]	0.0
(85,[0,2,9,10,11,...]	1.0
(85,[0,1,9,10,12,...]	1.0
(85,[0,1,9,10,11,...]	1.0
(85,[0,1,9,10,11,...]	1.0
(85,[0,5,9,10,11,...]	1.0
(85,[0,1,9,10,12,...]	0.0
(85,[0,1,9,10,12,...]	0.0
(85,[0,1,9,10,11,...]	1.0
(85,[0,1,9,10,11,...]	0.0
(85,[0,2,9,10,12,...]	0.0
(85,[0,1,9,10,12,...]	0.0
(85,[0,1,9,10,11,...]	0.0
(85,[0,2,9,10,13,...]	1.0

only showing top 20 rows

```
In [41]: ## Split the dataset to training set and test set
data_set = df_clean.select(['feature', 'incomeIndex'])

train_set, test_set = data_set.randomSplit([0.7, 0.3], 100)
```

(3). Modelling

```
In [42]: from pyspark.ml.classification import LogisticRegression
```

```
In [43]: ## Train the data and access the summary
lr = LogisticRegression(featuresCol = 'feature', labelCol = 'incomeIndex').fit(train_set)

training_summary = lr.summary
```

```
In [44]: training_summary.roc.show()
```


FPR	TPR
0.0	0.0
5.779010633379565E-5	0.04644018792916516
3.467406380027739...	0.08655583664618721
0.001098012020342...	0.12522587640043367
0.002369394359685622	0.16226960607155763
0.0039875173370319	0.19822912902059991
0.006125751271382339	0.2325623418865197
0.007917244567730004	0.2679797614745211
0.01103791030975497	0.2992410552945428
0.01410078594544614	0.33068305023491146
0.01716366158113731	0.3621250451752801
0.020804438280166437	0.39176002891217926
0.025369856680536294	0.4185037947235273
0.03005085529357374	0.44488615829418143
0.03467406380027739	0.4714492229851825
0.03981738326398521	0.4963859775930611
0.04455617198335645	0.5225876400433682
0.05010402219140083	0.5462594868088182
0.05640314378178456	0.5675822190097579
0.062471104946833104	0.5896277556920853

only showing top 20 rows

```
In [45]: train_FPR = []
for each in training_summary.roc.select('FPR').collect():
    train_FPR.append(each['FPR'])

train_TPR = []
for each in training_summary.roc.select('TPR').collect():
    train_TPR.append(each['TPR'])
```

```
In [46]: plt.plot(train_FPR, train_TPR)
plt.title('Train ROC Curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.show()
```



```
In [47]: training_summary.pr.show()
```

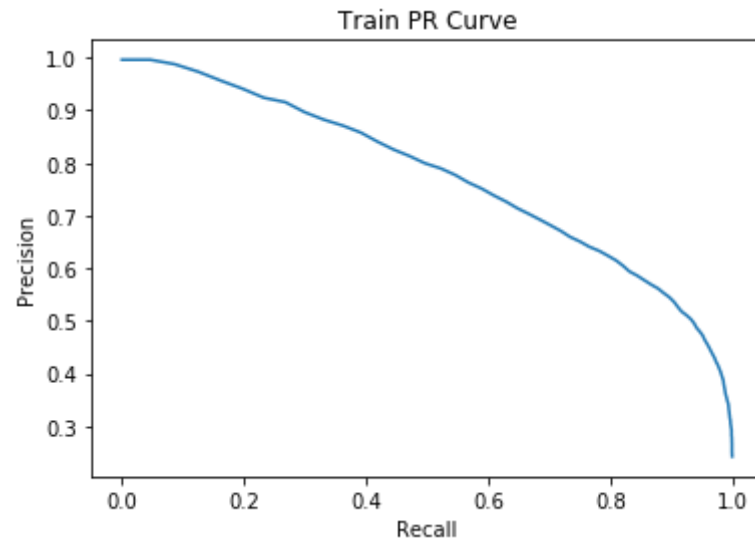
recall	precision
0.0	0.9961240310077519
0.04644018792916516	0.9961240310077519
0.08655583664618721	0.9876288659793815
0.12522587640043367	0.973314606741573
0.16226960607155763	0.9563365282215123
0.19822912902059991	0.9408233276157805
0.2325623418865197	0.923905240488155
0.2679797614745211	0.9154320987654321
0.2992410552945428	0.8965890633459664
0.33068305023491146	0.8823529411764706
0.3621250451752801	0.8709256844850065
0.39176002891217926	0.8575949367088608
0.4185037947235273	0.8406533575317604
0.44488615829418143	0.8256203890006707
0.4714492229851825	0.8130258647553755
0.4963859775930611	0.7994761350407451
0.5225876400433682	0.7895167895167895
0.5462594868088182	0.777120822622108
0.5675822190097579	0.7629341753704153
0.5896277556920853	0.7511510128913443

only showing top 20 rows

```
In [48]: train_recall = []
for each in training_summary.pr.select('recall').collect():
    train_recall.append(each['recall'])

train_precision = []
for each in training_summary.pr.select('precision').collect():
    train_precision.append(each['precision'])
```

```
In [49]: plt.plot(train_recall, train_precision)
plt.title('Train PR Curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.show()
```



(4). Evaluation

```
In [50]: test_summary = lr.evaluate(test_set)
```

```
In [51]: print 'The value of area under ROC is ', test_summary.areaUnderROC
```

The value of area under ROC is 0.902686157789

The value of area under ROC is 0.902686157789.

```
In [52]: print 'The accuracy is ', test_summary.accuracy
```

The accuracy is 0.849737735267

The accuracy is 0.849737735267 (around 85%) .

```
In [53]: predictionAndLabels = test_summary.predictions
```

```
In [54]: label_and_prediction = test_summary.predictions.select(['incomeIndex', 'prediction'])
```

```
In [55]: TP = 0
FP = 0
FN = 0
TN = 0

for each_lab_pred in label_and_prediction.collect():
    if each_lab_pred[0] == 1.0 and each_lab_pred[1] == 1.0:
        TP += 1
    elif each_lab_pred[0] == 0.0 and each_lab_pred[1] == 1.0:
        FN += 1
    elif each_lab_pred[0] == 1.0 and each_lab_pred[1] == 0.0:
        FP += 1
    elif each_lab_pred[0] == 0.0 and each_lab_pred[1] == 0.0:
        TN += 1
```

```
In [56]: confusion_matrix = np.array([[TP, FP], [TN, FN]])
```

```
In [57]: print 'The confusion matrix is \n', confusion_matrix
```

The confusion matrix is
[[1362 945]
 [6900 516]]

```

In [58]: import itertools

state_label = ['Positive', 'Negative']
predi_label = ['True', 'False']

plt.title('Confusion Matrix')
plt.imshow(confusion_matrix, interpolation = 'nearest', cmap = 'Blues')
plt.colorbar()
plt.xticks(np.arange(2), predi_label)
plt.yticks(np.arange(2), state_label)

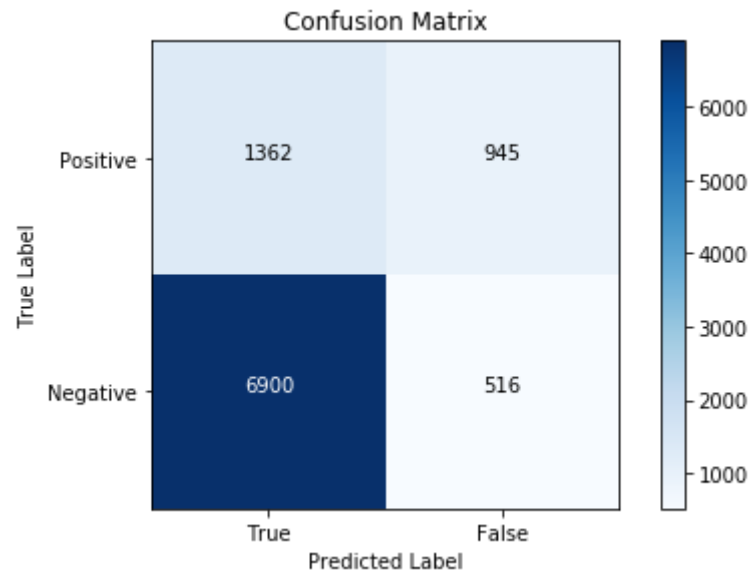
for i, j in itertools.product(range(confusion_matrix.shape[0]), range(confusion_matrix.shape[1])):
    plt.text(j, i, format(confusion_matrix[i, j], 'd'),
             horizontalalignment = 'center', color = 'white' if confusion_matrix[i, j] > confusion_matri

plt.tight_layout()

plt.xlabel('Predicted Label')
plt.ylabel('True Label')

plt.show()

```



In []: